




# Pair-Programming con IA para conseguir código más eficiente en .NET





 jcanton

 jcant0n

## Javier Cantón Ferrero

Javier is a Computer Science Engineer who has always had a passion for 3D graphics and software architecture. His professional achievements include being MVP for Windows DirectX and DirectX XNA for nine years, Xbox Ambassador, as well as Microsoft Student Partner and Microsoft Most Valuable Student during his years at college. Currently he works as Chief Technology Innovation Officer.






Evergine is our [cross-platform graphics engine designed for corporate applications](#). It has been created to offer enterprise solutions with 3D, Augmented, Virtual and Mixed Reality capabilities.

A flexible and versatile tool to create unique visual experiences, adapted to the needs and functionalities of all types of industries.

Evergine is [free to use](#), adapts to any platform and is ready to be multi-platform.



Windows



Android



Linux



MacOs + iOS

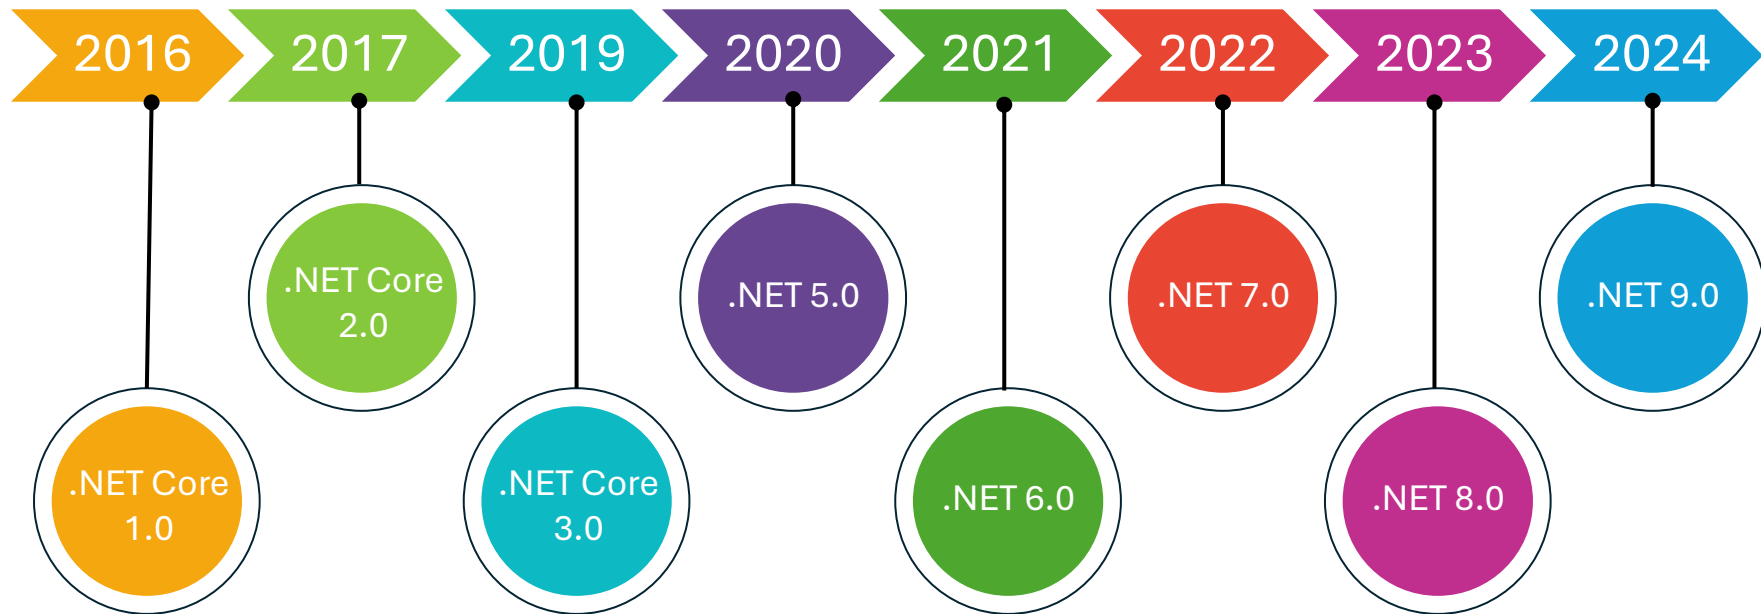


Web

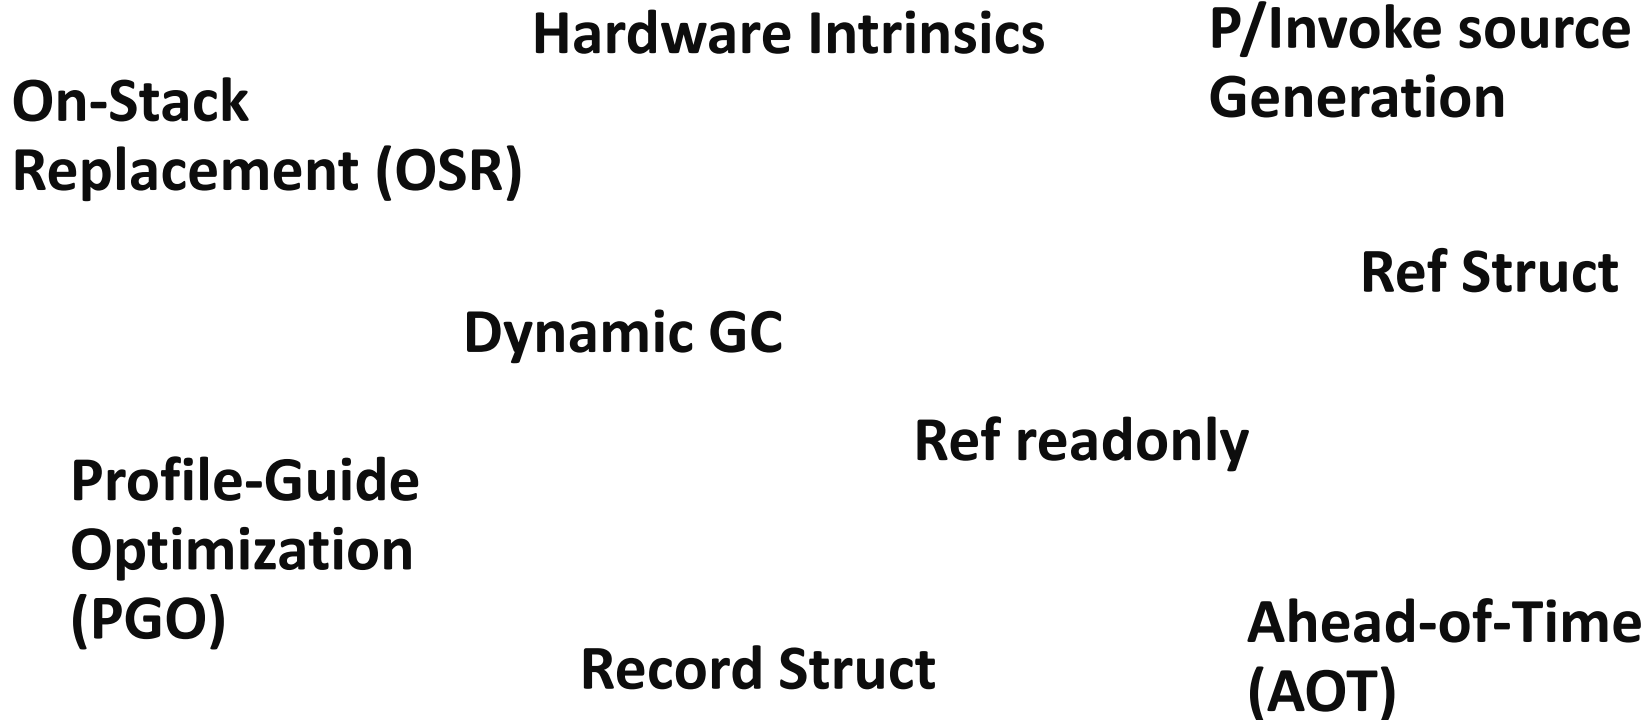


# **.NET Performance**

# .NET Evolution



# NET features



# Micro-Optimization

---

- BCL and Runtime
- Real-Time applications
- Graphics development
- ML or Math libraries





# What is faster (Assignment)?

```
void A(int x, int y) {
    for (int i = 0; i < 10000; i++)
    {
        a[i] += i;
        a[i] += x;
        a[i] += y;
    }
}
```



```
void B(int x, int y) {
    for (int i = 0; i < 10000; i++)
    {
        a[i] = a[i] + i;
        a[i] = a[i] + x;
        a[i] = a[i] + y;
    }
}
```

Method	x	y	Mean	Error	StdDev	Ratio
A	10	10	6.938 us	0.0237 us	0.0221 us	1.00
B	10	10	8.489 us	0.0072 us	0.0067 us	1.22

# What is faster (Nullable)?

```
int A(int? x)
{
    for (int i = 0; i < 1000; i++)
        x++;
    return x.Value;
}
```

```
int B(int x)
{
    for (int i = 0; i < 1000; i++)
        x++;
    return x;
}
```



Method	x	Mean	Error	StdDev	Ratio
A	0	615.5 ns	1.99 ns	1.77 ns	1.00
B	0	264.7 ns	3.63 ns	3.40 ns	0.43

# What is faster (Collection expressions)?

```
void A() {
    List<int> a =
        [1, 2, 3, 5];
}
```

```
void B() {
    List<int> b = new List<int>
        {1, 2, 3, 5};
}
```



Method	Mean	Error	StdDev	Median
A	8.741 ns	0.2440 ns	0.7157 ns	8.541 ns
B	13.399 ns	0.2541 ns	0.3024 ns	13.410 ns

# What is faster (Try-catch)?

```
int A(int x) {
    try {
        for(int i = 0; i < 1000; i++)
            x++;
    }
    catch { }
    return x;
}
```

```
public int B(int x) {
    try {
        for(int i = 0; i < 1000; i++)
            x++;
    }
    catch { throw; }
    return x;
}
```




Method	x	Mean	Error	StdDev	Median	Ratio
A	0	1,782.5 ns	35.35 ns	53.98 ns	1,751.2 ns	1.00
B	0	361.9 ns	8.14 ns	24.01 ns	358.4 ns	0.21

# Measure

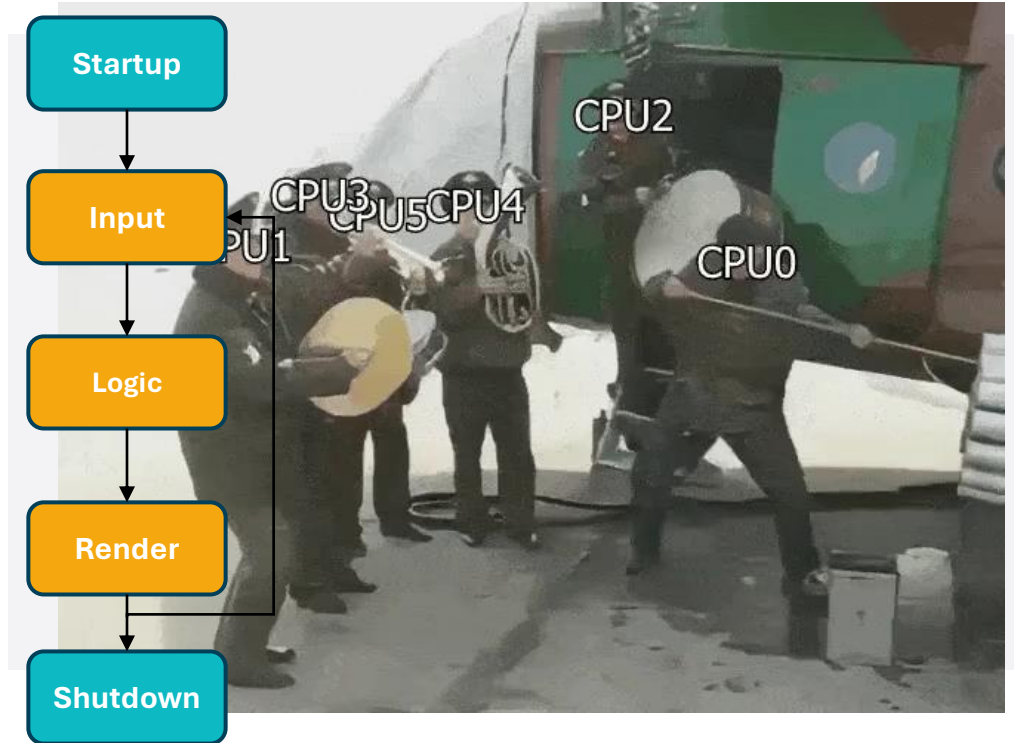


**BenchmarkDotNet**  
Powerful .NET library for benchmarking

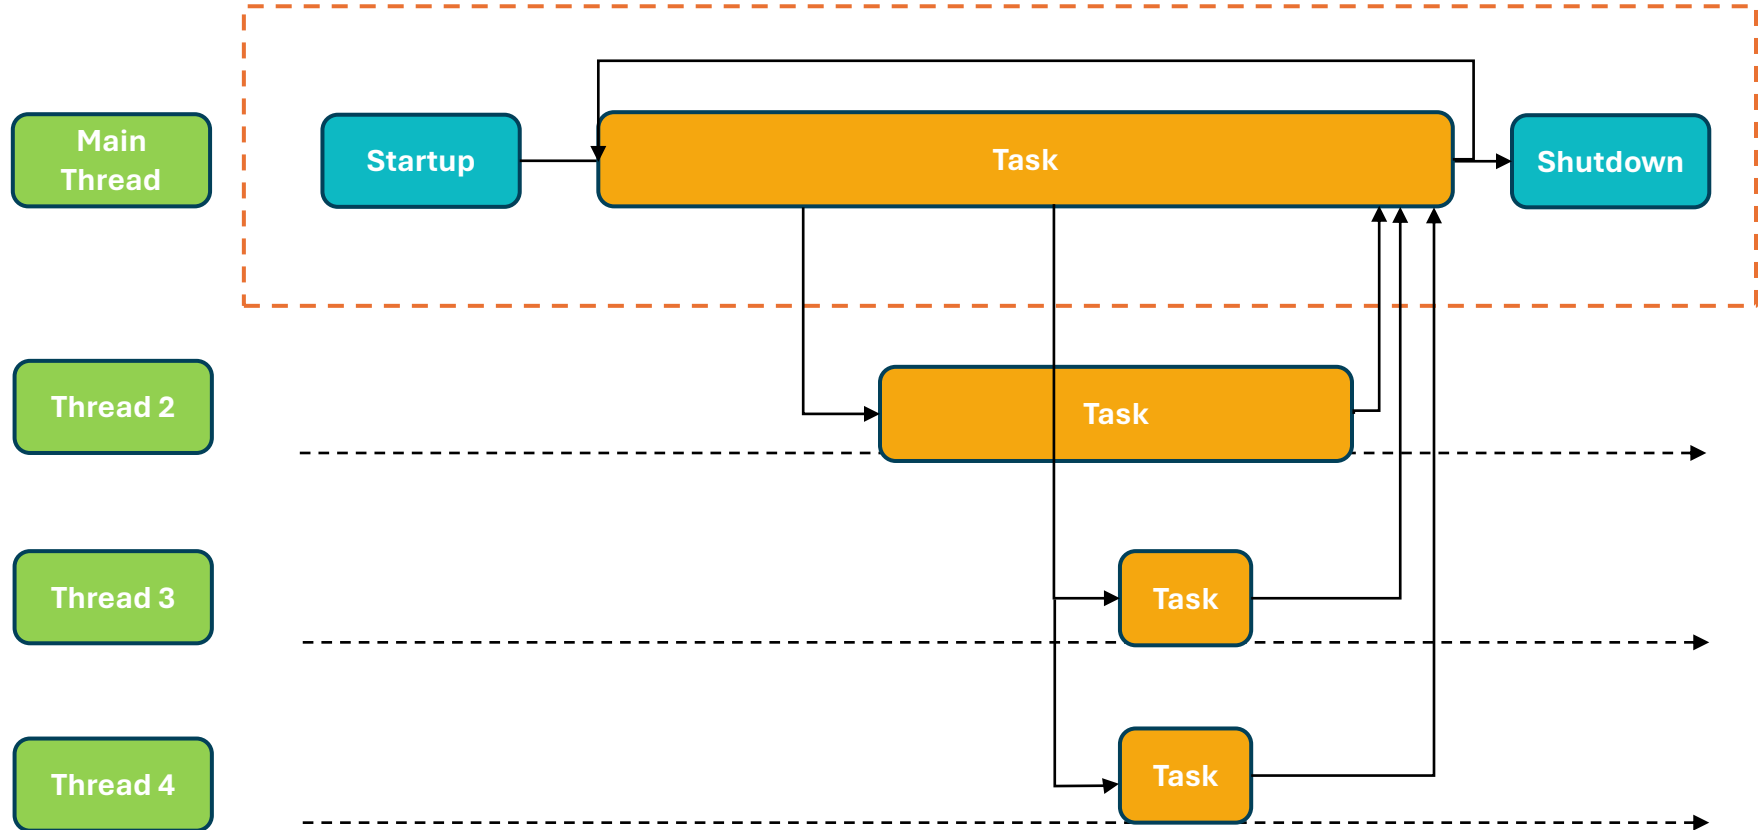


**JobSystem**

# Usually 3D app



# JobSystem





# GenAI for coding


# News

The New York Times

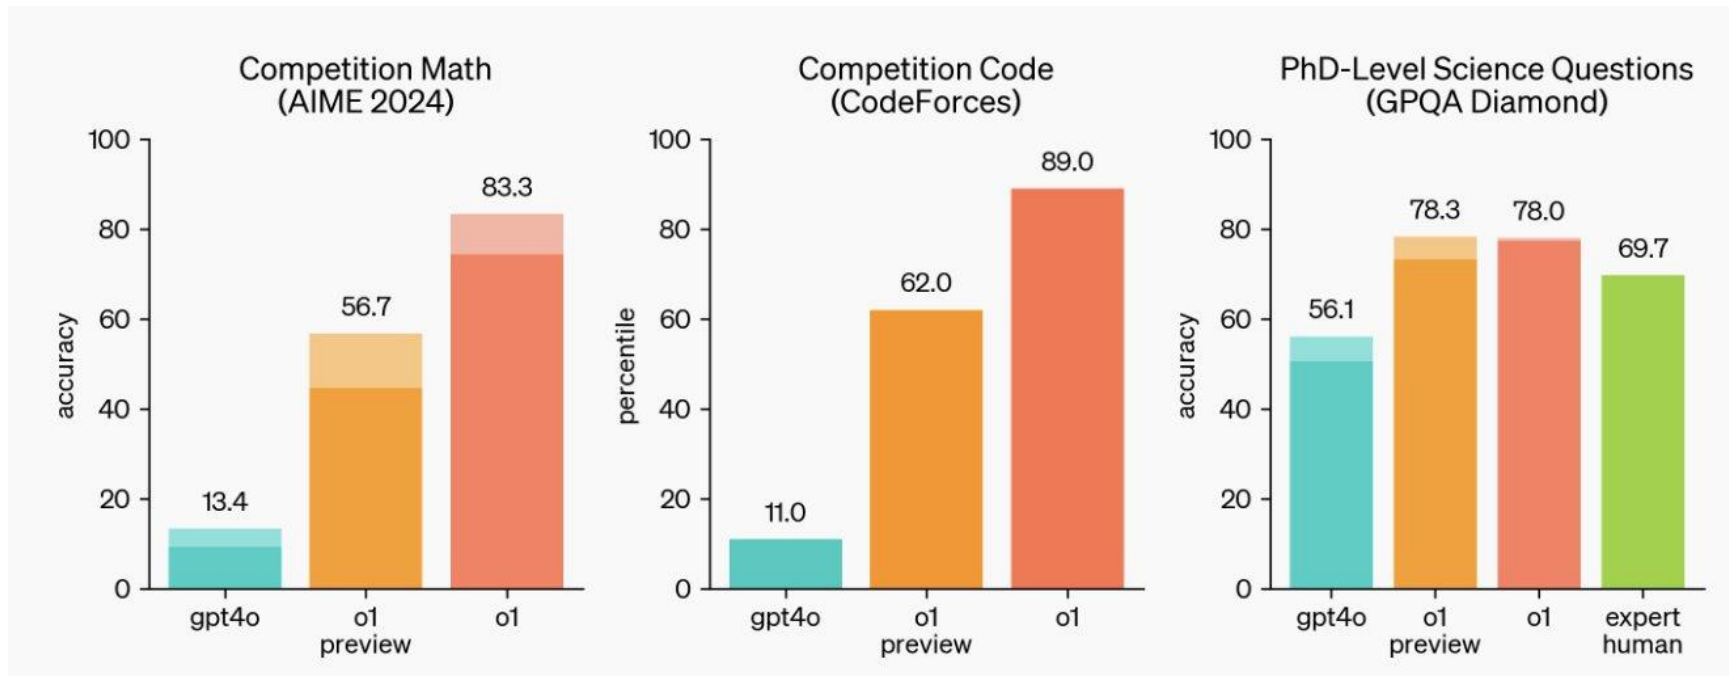
OPINION  
FARHAD MANJOO

## It's the End of Computer Programming as We Know It. (And I Feel Fine.)

June 2, 2023



[https://twitter.com/gadget\\_ry](https://twitter.com/gadget_ry)  
@gadget\_ry, Developer, Development Tools, Open Source,



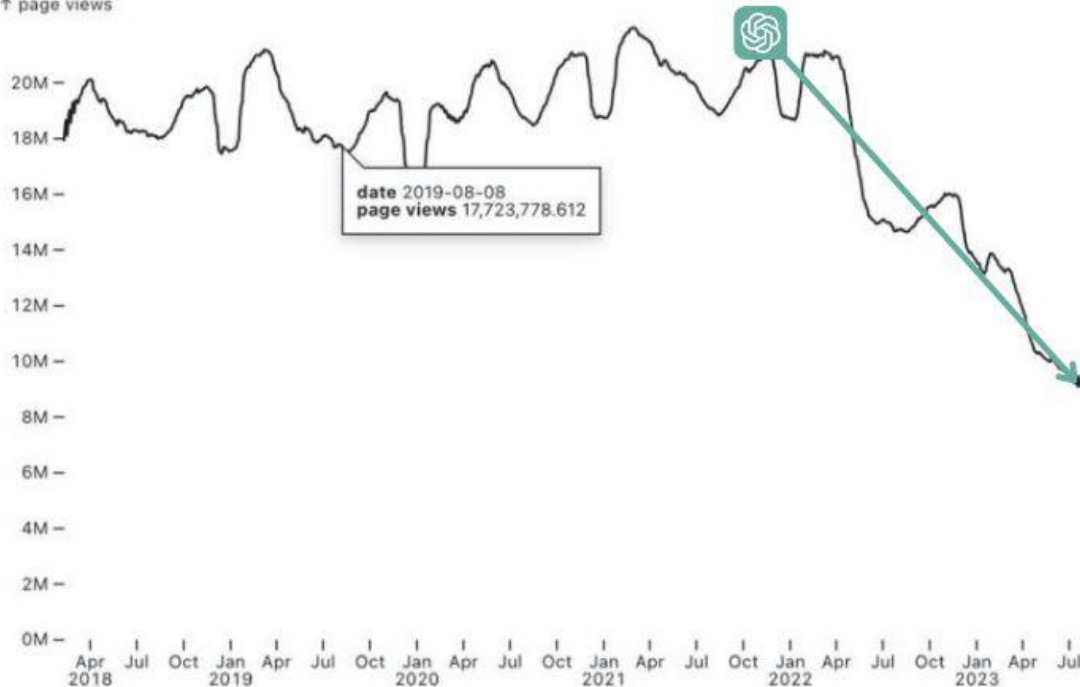
# stackoverflow website traffic

Select one series  
for traffic

page views

↑ page views

ChatGPT released



# LLMs for coding




Chatgpt 4o

OpenAI o1 preview



V0 vercel



Claude 3.5



Github Copilot

# Pair-Programming

# Pair-Programming

- **Driver**
  - The person who write the code.
  - Handles the implementation details, syntax and immediate problem-solving.
- **Navigator**
  - The person who reviews and provides guidance.
  - Thinks about the overall approach



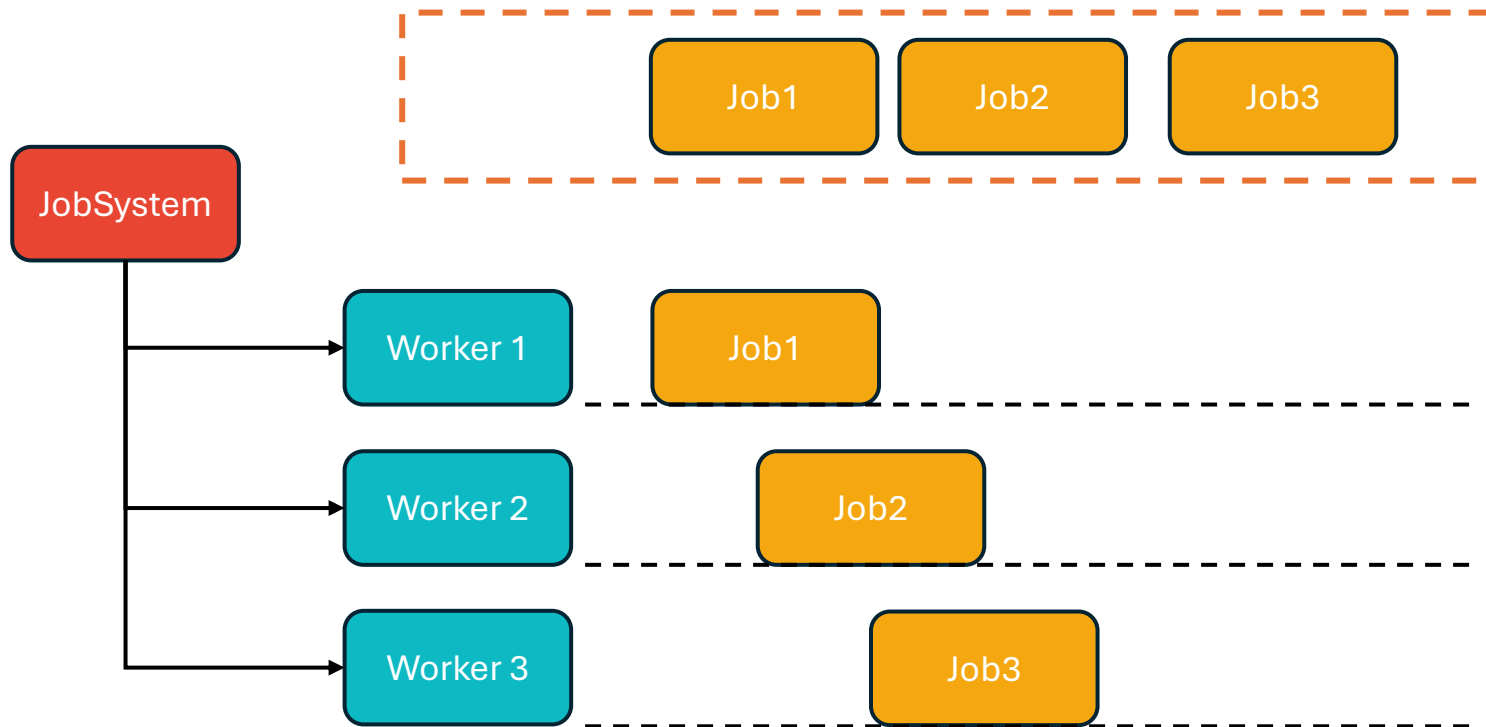
# Rubber Duck Programming

- **How it works:**
  - Pace your “duck” in front of you.
  - Explain your code or problem you are facing as if the duck needs to understand it.
  - As you verbalize your logic, your brain often identifies gaps or bugs that were previously overlooked.





# Basic JobSystem



# JobSystem

```
namespace JobSystemTest
{
    5 references | jcant0n, 71 days ago | 1 author, 1 change
    public class JobSystem : IDisposable
    {
        public Worker[] Workers;
        public readonly uint NumThreads;
        private uint nextQueueIndex;
        private bool isRunning;
        private bool disposed;

        6 references | jcant0n, 71 days ago | 1 author, 1 change
        public class Context...

        1 reference | jcant0n, 71 days ago | 1 author, 1 change
        public JobSystem(uint maxThreadCount = 0)...

        1 reference | jcant0n, 71 days ago | 1 author, 1 change
        public void Dispose()...

        1 reference | jcant0n, 71 days ago | 1 author, 1 change
        public bool IsBusy(Context context)...

        1 reference | jcant0n, 71 days ago | 1 author, 1 change
        public void Wait(Context context)...

        0 references | jcant0n, 71 days ago | 1 author, 1 change
        public void Execute(Context context, Action<JobArgs> function)...

        1 reference | jcant0n, 71 days ago | 1 author, 1 change
        public void Dispatch(Context content, uint jobCount, uint groupSize, Action<JobArgs> function)...
    }
}
```

# Execute

0 references | jcant0n, 71 days ago | 1 author, 1 change

```
public void Execute(Context context, Action<JobArgs> function)
{
    Interlocked.Increment(ref context.PendingJobs);

    Job job = new Job
    {
        Function = function,
        Context = context,
        GroupID = 0,
        GroupJobOffset = 0,
        GroupJobEnd = 1,
    };

    uint queueIndex = Interlocked.Increment(ref nextQueueIndex) % NumThreads;
    Worker w = Workers[queueIndex];
    w.JobQueue.Enqueue(job);
    w.Signal.Set();
}
```

# Benchmark (AI)

```
[Benchmark]
public void MultiplyMatrixWithParallelFor()
{
    // Multiply matrices using Parallel.For
    Parallel.For(0, matrixSize, i =>
    {
        for (int j = 0; j < matrixSize; j++)
        {
            int sum = 0;
            for (int k = 0; k < matrixSize; k++)
            {
                sum += matrix1[i, k] * matrix2[k, j];
            }
            result[i, j] = sum;
        }
    });
}
```

```
[Benchmark]
public void MultiplyMatrixWithJobSystemDispatch()
{
    // Create a new context for this benchmark
    var context = new JobsContext();

    uint jobCount = (uint)matrixSize; // One job per row
    uint groupSize = 10; // Adjust group size as needed

    // Use JobSystem.Dispatch to distribute the work across jobs
    jobSystem.Dispatch(context, jobCount, groupSize, (args) =>
    {
        int i = (int)args.JobIndex;
        for (int j = 0; j < matrixSize; j++)
        {
            int sum = 0;
            for (int k = 0; k < matrixSize; k++)
            {
                sum += matrix1[i, k] * matrix2[k, j];
            }
            result[i, j] = sum;
        }
    });
}
```

# New ideas

## 2. 3. Optimize JobSystem Internals:

- **Reduce Synchronization Overhead:**
  - **Avoid Excessive Locking:** Check if there are any unnecessary locks or synchronization primitives in your `JobSystem`. Even though you're using `ConcurrentQueue` and `Interlocked` operations, ensure that there's no contention.
  - **Use SpinWait:** In your worker threads, consider using `SpinWait` instead of `ManualResetEventSlim` for waiting. This can reduce context switching overhead.

```
csharp
Copy code

var spinWait = new SpinWait();
while (isRunning)
{
    if (QueuePerWorker[threadID].TryDequeue(out var job))
    {
        job.Execute();
    }
    else
    {
        // Attempt to steal work or spin
        // Stealing Logic here
        spinWait.SpinOnce();
    }
}
```

# More

I have  
implem



The

I ur

you

no

To

on

lor


## B. Use Structs and Ref Structs Wisely

Using `struct` s can reduce heap allocations, but large structs can cause performance issues due to copying.

Solution: Keep `struct` s Small and Use `in` Parameters

Ensure your structs are small and consider passing them by reference.

Example:

 Copy code

csharp

```
public readonly struct JobArgs
{
    public readonly uint JobIndex;
    public readonly uint GroupID;
    public readonly uint GroupIndex;

    public JobArgs(uint jobIndex, uint groupID, uint groupIndex)
    {
        JobIndex = jobIndex;
        GroupID = groupID;
        GroupIndex = groupIndex;
    }
}
```

de

e in

ly

# Inmutable struct

```
/// <summary>
/// Represents the arguments passed to a job function.
/// </summary>
[StructLayout(LayoutKind.Explicit, Size = 16)]
7 references | jcant0n, 66 days ago | 1 author, 4 changes
public readonly struct JobArgs
{
    /// <summary>
    /// The index of the current job.
    /// </summary>
    [FieldOffset(0)]
    public readonly uint JobIndex;

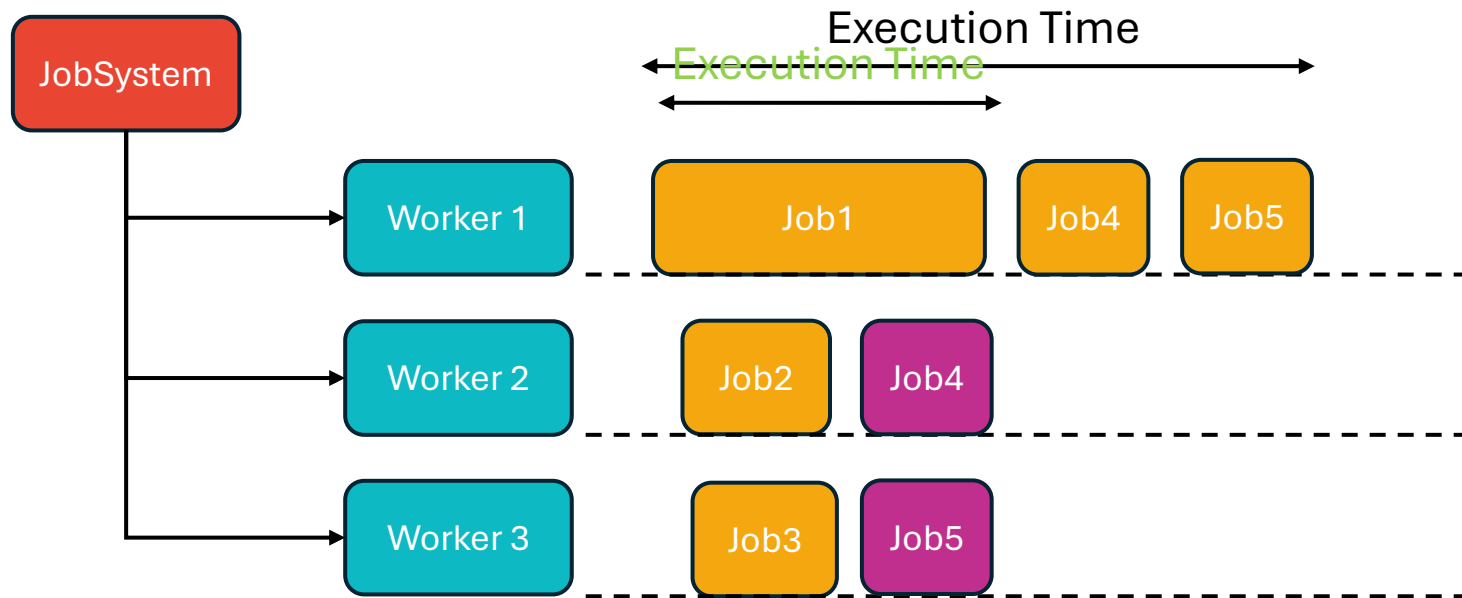
    /// <summary>
    /// The ID of the group this job belongs to.
    /// </summary>
    [FieldOffset(4)]
    public readonly uint GroupID;

    /// <summary>
    /// The index of this job within its group.
    /// </summary>
    [FieldOffset(8)]
    public readonly uint GroupIndex;

    /// <summary>
    /// Padding to ensure 16-byte alignment.
    /// </summary>
    [FieldOffset(12)]
    private readonly uint padding;
}
```



# Stealing jobs





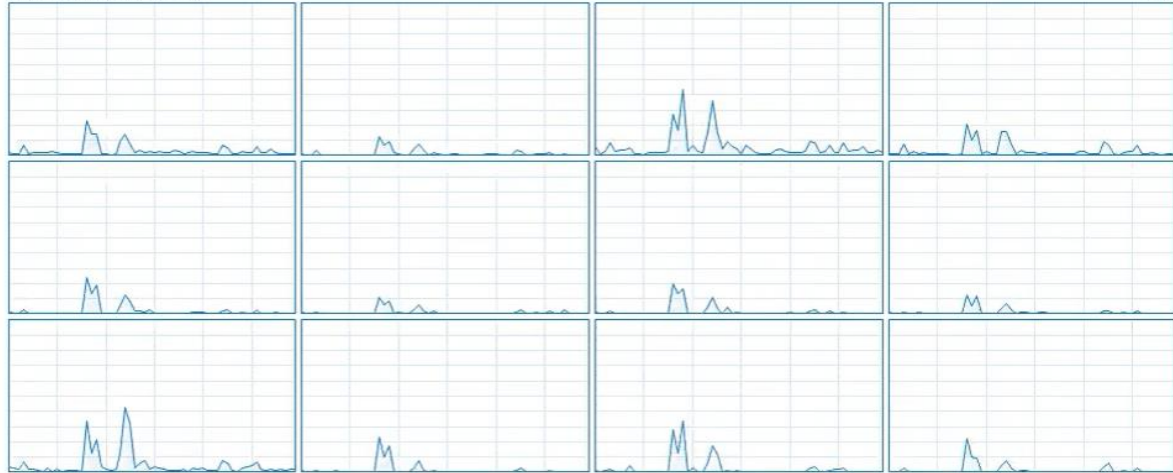
# CPU usage

## CPU

AMD Ryzen 5 3600 6-Core Processor

Auslastung in 60 Sekunden (%)

100%



Auslastung	Geschwindigkeit	Basisgeschwindigkeit:	3.60 GHz	
1%	3.79 GHz	Sockets:	1	
		Kerne:	6	
Prozesse	Threads	Handles	Logische Prozessoren:	12
174	2231	72014	Virtualisierung:	Deaktiviert
Betriebszeit			Hyper-V-Unterstützung:	Ja
0:02:54:23			L1-Cache:	384 KB
			L2-Cache:	3.0 MB
			L3-Cache:	32.0 MB

[monitor öffnen](#)



# Random

[MethodImpl(MethodImplOptions.AggressiveInlining)]

2 references | jcant0n, 62 days ago | 1 author, 2 changes

public ulong NextUInt64()

```
{  
    ulong x = state;  
    x ^= x << 13;  
    x ^= x >> 7;  
    x ^= x << 17;  
    state = x;  
    return x;  
}
```



public Xoshiro256StarStar(ulong seed)

```
{  
    if (seed == 0)  
        seed = 0xdeadbeef;  
  
    state0 = seed;  
    state1 = SplitMix64(seed + 1);  
    state2 = SplitMix64(seed + 2);  
    state3 = SplitMix64(seed + 3);  
}
```

/// <summary> Implements the SplitMix64 algorithm, used for

[MethodImpl(MethodImplOptions.AggressiveInlining)]

3 references | jcant0n, 66 days ago | 1 author, 2 changes

private static ulong SplitMix64(ulong x)

```
{  
    x += 0x9e3779b97f4a7c15;  
    x = (x ^ (x >> 30)) * 0xbf58476d1ce4e5b9;  
    x = (x ^ (x >> 27)) * 0x94d049bb133111eb;  
    return x ^ (x >> 31);  
}
```

Job=DefaultJob

Method	Mean	Allocated
SystemRandom_Next	6.137 ms	3 B
XorShiftRandom_Next	2.032 ms	2 B
Xoshiro256Random_Next	1.856 ms	1 B



# Final Performance



```
BenchmarkDotNet v0.14.0, Windows 11 (10.0.22631.4037/23H2/2023Update/SunValley3)
13th Gen Intel Core i7-13700KF, 1 CPU, 24 logical and 16 physical cores
.NET SDK 8.0.202
```

```
[Host] : .NET 8.0.3 (8.0.324.11423), X64 RyuJIT AVX2
```

```
Job-IOEUFW : .NET 8.0.3 (8.0.324.11423), X64 RyuJIT AVX2
```

```
Job=Job-IOEUFW InvocationCount=1 UnrollFactor=1
```

```
RatioSD=0.02
```

Method	Mean	Ratio	Allocated	Alloc Ratio
MultiplyMatrixWithParallelFor	14.28 ms	1.00	7248 B	1.00
MultiplyMatrixWithJobSystemDispatch	11.08 ms	<u>0.78</u>	760 B	<u>0.10</u>

22% más rápido y un 90% menos Alloc

# Conclusions

# Pros and Cons

## POSITIVES



Speed up coding

Enhance Code Naming

Helpful in brainstorming solutions

Great to generate performance tests

Excellent for language translation



## NEGATIVES

Limited Context Understanding

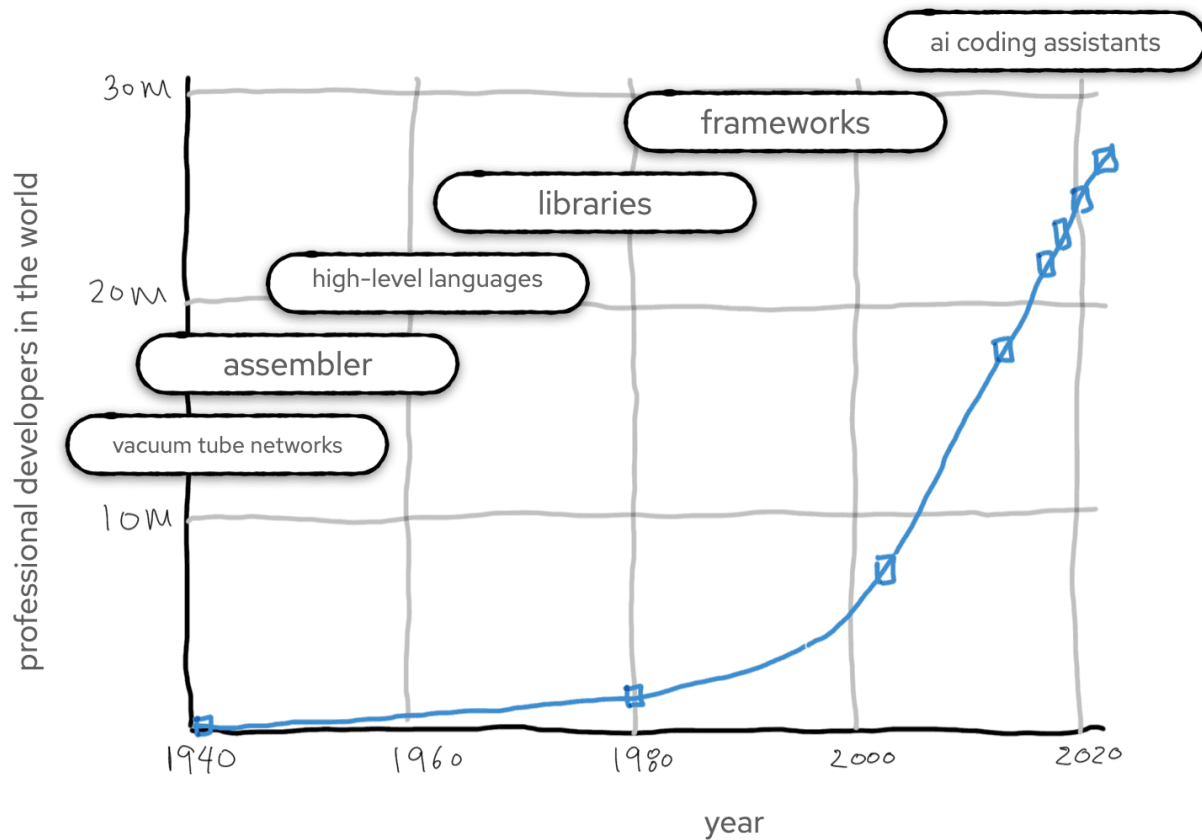
It could saturate users

Overconfidence in AI


Data security risks

Overly generic default solutions

# Productivity Tool



# Questions





EL EVENTO SOBRE

TECNOLOGÍAS  
**CLOUD, WEB  
Y DATA**

